

---

# Multimodal AI: Towards Deterministic Verification Layers for Large Language Models

---

**Author:**

**Nikhil Tripathi**

Solution Architect with over a decade of experience in enterprise IT architecture, LLMs, AI systems integration, and digital transformation.

**Email:** *[nikhil.k.tripathi@gmail.com](mailto:nikhil.k.tripathi@gmail.com)*

---

# 1. Executive Summary

Large Language Models (LLMs) can now interpret and generate text, images, and code fluently — but their reasoning remains **probabilistic**, not **proven**. Humans verify through **perception, logic, and experience**; LLMs predict words based on **statistical likelihood**.

To close this gap, we introduce the **Deterministic Verification Layer (DVL)** — a symbolic, rule-based layer that **computes, validates, and explains** before the LLM is allowed to narrate.

The DVL fuses:

- **RDF** knowledge graphs for structured semantics
- **SHACL** constraint reasoning for logical validation
- **Deterministic computation** for proof-based accuracy
- **Provenance tracking** for auditability
- A **policy of refusal** for impossible statements

This approach gives machines a **human-like verification sense** — they “check before believing.”

## Current Status

- **Math Verification Layer — Completed (PoC ready)**  
Deterministic evaluator + RDF/SHACL validation + Gradio UI in Google Colab.
  - **Biology Validation Layer — In Progress**  
Ontology (Plant/Animal) and exception rules implemented; controller wiring and test coverage underway.
-

---

## 2. Problem Context

### 2.1 The Human–Machine Perception Gap

Humans assess reality through multimodal perception — sight, sound, memory, and logic — ensuring internal consistency before acting.

LLMs, however, operate on **token prediction**, not factual verification.

They can:

- Produce numerically wrong yet fluent answers,
- Contradict themselves between sessions,
- Lack an audit trail or ground truth link.

When deployed in regulated sectors (finance, law, engineering, healthcare), these weaknesses become structural risks.

### 2.2 Why Existing Safeguards Fall Short

Approach	Strength	Limitation
Guardrails & Filters	Style & safety	Don't verify correctness
RLHF	Improves helpfulness	Doesn't prove truth
Confidence Scores	Internal metrics	Non-reproducible
Provenance Frameworks (C2PA)	Track content origin	Don't test factuality

Hence, a **verification-first** architecture — independent of model probability — is needed.

---

## 3. The Deterministic Verification Layer (DVL)

### 3.1 Concept Overview

The DVL acts as a **semantic firewall** between the LLM and its users.

Each claim generated by an AI system is:

1. Encoded into **RDF triples** (facts),
2. Validated by **SHACL constraints**,
3. Recomputed or verified deterministically,
4. Returned **only if validated** — else politely refused.

### 3.2 The Human-Perception Analogy

Human Function	DVL Equivalent
Sensory Perception	Multimodal input parsing (text, image, audio)
Logical Reasoning	RDF graph reasoning
Consistency Checking	SHACL constraint validation
Calculation & Estimation	Deterministic computation engine
Moral/Contextual Response	Polite refusal on impossibilities

The DVL brings *human-like perception and judgment* to AI — ensuring it reasons before responding.

## 4. Modular Architecture

### 4.1 Layered Overview

-----+   USER / FRONTEND INTERFACE   (Gradio UI / Web App / Teams Bot / MCP Client)  -----+   DETERMINISTIC VERIFICATION LAYER (DVL)  -----+   1. Input Interpreter   Entity & Relation Extraction   2. RDF Encoder   Structured Triple Generation   3. Domain Validator   SHACL Constraint Validation   4. Deterministic Engine   Math / Logic Verification   5. Exception Reasoner   Impossible / Context Rules   6. Provenance Recorder   Hash, Timestamp, Audit Trail   7. Language Narrator (LLM)   Post-Validation Explanation  -----+   FOUNDATION MODEL (DeepSeek / GPT / Claude)   Optional – used for narration after validation  -----+   TOOL ORCHESTRATION LAYER (MCP-Compatible)   math.evaluate   rdf.query   shacl.validate   provenance.log    -----+  -----+	
--	--

### 4.2 Module Descriptions

#### 1. Input Interpreter

- Accepts text or multimodal inputs (e.g., image + caption).
- Detects entities, verbs, and relationships.
- Outputs structured claims for validation, e.g.:
  - “The cat ran.” → (Cat, performed, Running)
  - “5 + 7 = 12.” → (Equation, operand1, operand2, operator).

## 2. RDF Encoder

- Encodes inputs as RDF triples with namespaces and unique identifiers.
- Example:
- `ex:RunEvent1 a act:Running ;`
- `act:agent ex:ThisCat ;`
- `act:time "2025-10-26T12:00Z"^^xsd:dateTime .`

## 3. Domain Validator (SHACL)

- Applies SHACL shapes to enforce rules:
  - Structural (e.g., every event must have an agent).
  - Semantic (e.g., only animals can run).
  - SPARQL constraints for exception messages.
- Returns deterministic validation reports.

## 4. Deterministic Engine

- Recomputes mathematical or logical expressions exactly.
- Implements safe AST parsing using Python's `decimal` for precision.
- Used in the Math PoC for verifiable arithmetic and Boolean logic.

## 5. Exception Reasoner

- Contains explicit “impossible” rules (e.g., `Plant cannotPerform Running`).
- When violated, returns:

“I politely refuse your statement. What you are saying is highly unlikely unless there is some exception.”

## 6. Provenance Recorder

- Captures rule version, timestamp, and validation result.
- Generates immutable logs suitable for audit or blockchain anchoring.
- Aligns with **C2PA** and **ISO 42001** provenance standards.

## 7. Language Narrator

- Invokes the LLM **only after validation passes**.
  - Generates natural, human-readable explanations of validated facts.
  - For invalid claims, produces the standardized **polite refusal** message.
-

## 4.3 Tool Orchestration (MCP Integration)

All components can be exposed as MCP-style tools:

<code>math.evaluate()</code>	→ deterministic calculator
<code>rdf.query()</code>	→ triple retrieval
<code>shacl.validate()</code>	→ constraint checking
<code>provenance.log()</code>	→ audit record creation

This allows modular, service-based scaling while ensuring strict sequencing — validation **always precedes** narration.




---

## 5. Proofs of Concept

### 5.1 Mathematical Verification PoC — Status: Completed

**Stack:** Python + RDFLib + PySHACL + Gradio

**Functionality:**

1. Parse arithmetic expression or DeepSeek JSON response.
2. Compute deterministically with high precision.
3. Encode as RDF triples (`math:operand1`, `math:operator`, `math:computedResult`).
4. Validate structure and datatypes with SHACL.
5. Return a structured verdict:  Valid /  Mismatch /  Unverified.
6. Log timestamped provenance.

**UI:** Interactive Gradio app (Google Colab).

**Optional:** DeepSeek API integration for automatic JSON extraction.

#### Output Example

```
{
  "status": "valid",
  "answer": {"computed": "13.0000"},
  "evidence": {"constraints_passed": true},
  "provenance": {"timestamp_utc": "2025-10-26T12:34Z"}
}
```

---

### 5.2 Biological Validation PoC — Status: In Progress

**Domain:** Biology — Plants & Animals





**Ontology:** Plant, Animal, Cat, Photosynthesis, Running

**Implemented:**

- RDF ontology and capabilities (`Plant canPerform Photosynthesis`, `Cat canPerform Running`)

- Exception rules (Plant cannotPerform Running, Cat cannotPerform Photosynthesis)
- SHACL shapes + SPARQL constraints to enforce capability logic
- **Polite refusal mechanism** triggered on biologically impossible statements

#### Pending Work:

- Controller wiring (capability → event → SHACL → response)
- Test cases (cat ran , plant ran , plant photosynthesized , cat photosynthesized )
- Integration with MCP tool endpoints (rdf.query, shacl.validate)

#### Expected Outcome:

Validation results identical to deterministic perception — only feasible actions accepted.

---

## 6. Expanding Toward Full Multimodal Perception

### 6.1 Domain Modules

- **Mathematics:** Arithmetic, algebraic consistency, logical evaluation.
- **Biology:** Capability mapping, ecological logic.
- **Languages:** Lemma → Action mapping, tense and polarity handling.
- **Engineering:** Mechanical feasibility, energy constraints, material logic.
- **Social Sciences:** Normative cause-effect validation.
- **Arts:** Conceptual relations and emotional semantics.

Each module is stored as an RDF graph pack:

/core, /capabilities, /events, /shapes, /exceptions.

### 6.2 Bridge Graphs

Domain bridges link:

- Language ↔ Action (ran → act:Running)
- Vision ↔ Object (detected image class → RDF entity)
- Math ↔ Physics (quantitative laws → physical validation)

These enable human-like **cross-domain perception** across text, image, and numeric modalities.

---

## 7. Evaluation Metrics

Metric	Description	Target
Deterministic Accuracy	Correctness after validation	≥ 95%
Explainability	Outputs with RDF + rationale	100%
Latency Overhead	Added verification time	< 500 ms
Consistency	Reproducibility across runs	100%
Provenance Integrity	Immutable audit chain	Full traceability

## 8. Strategic Significance

- Establishes a **trust layer** for all multimodal AI systems.
- Enables **verifiable intelligence** — beyond token prediction.
- Provides **regulatory compliance** under EU AI Act / NIST RMF / ISO 42001.
- Bridges **symbolic precision** with **foundation model perception**.
- Lays groundwork for **Provenance-as-a-Service (PaaS)** ecosystems.

## 9. Roadmap

Phase	Focus	Deliverables	Status
M0	Math PoC	Deterministic RDF/SHACL validator + Gradio UI	✅ Completed
M1	Biology PoC	Ontology + Capabilities + SHACL Exceptions + Controller Tests	🔄 In Progress
M2	Language & Lemma Mapping	NLP → Ontology Bridge + Coreference Resolution	Planned
M3	MCP Orchestration	Tool Endpoints ( <code>math.evaluate</code> , <code>rdf.query</code> , <code>shacl.validate</code> , <code>provenance.log</code> )	Planned
M4	Domain Expansion	Arts / Engineering / SocialScience Packs + Bridge Graphs	Planned
M5	Provenance & Dashboard	Immutable Logs + Metrics Interface	Planned
M6	Publication & Demo	Public Repo + White Paper Release	Planned



---

## 10. Human-Centric Interpretation

LLMs were built *for* humans — to converse naturally — but not *like* humans, who verify facts before speaking.

By adding the DVL, we give LLMs a **verification sense**:

- They *see* through RDF structure,
- They *judge* through SHACL validation,
- They *re-check* via deterministic computation,
- They *respond* with contextual politeness when logic fails.

Thus, DVL is the missing bridge between **human perception** and **machine reasoning**.

---

## 11. Conclusion

The **Deterministic Verification Layer** transforms large language models into **auditable reasoning systems**.




It merges **symbolic precision**, **mathematical determinism**, and **human-like perception** into a single verifiable architecture.

By computing and validating before narrating, AI systems evolve from *probabilistic assistants* to *accountable intelligence partners* — the foundation of trustworthy multimodal AI.

---

## 12. Appendix

### Appendix A — Math Verification Layer

- **Stack:** Python + RDFLib + PySHACL + Gradio
- **Functions:** Deterministic evaluation → RDF encoding → SHACL validation → Provenance logging.
- **Verdicts:**  Valid /  Mismatch /  Unverified.
- **UI:** Gradio interface with shareable Colab link.

### Appendix B — Biology Validation Layer

**Status:**  In Progress

- RDF ontology + SHACL exception constraints implemented.
- Controller + test cases pending.
- Produces polite refusal on impossible events.
- Expands human-like plausibility perception in LLM systems.